

The Surprise Machine Principle: A Hypothesized Law-Like Pattern in Nonequilibrium Adaptive Systems

Keith Ballinger

August 18, 2025

Abstract

We hypothesize a statistical principle for the emergence and growth of *surprise machines*—systems that persistently generate outputs with high observer-relative surprise—in any nonequilibrium environment meeting three prerequisites: (i) sustained free-energy flux, (ii) memory and computational capacity, and (iii) adaptive subsystems engaged in mutual modeling. We formalize this as the Surprise Machine Principle (SMP), which consists of a provable thermodynamic-communication bound on the maximum rate of surprise and a conjectured tendency for systems to evolve toward this bound. We prove the bound and then provide extensive empirical support for the conjectured tendency. Simulations across a wide range of models—including finite-state machines, chaotic maps, cellular automata, neural networks, hidden semi-Markov models, and probabilistic context-free grammars—demonstrate bound adherence, co-evolutionary arms races, and cost-sensitive escalation. A key experiment directly demonstrates that sophisticated observers create selection pressure for generators of higher computational complexity, providing a crucial link between the information-theoretic surprise and its thermodynamic cost. All code and data are available for full reproducibility.

1 Introduction

Many nonequilibrium systems—from predator–prey arms races to adversarial AI—generate sustained unpredictability relative to their observers. We call such systems *surprise machines*. We propose the ****Surprise Machine Principle (SMP)****, a hypothesized law-like pattern consisting of two core components:

1. **A Provable Bound:** A formal thermodynamic-communication envelope that limits the maximum rate of surprise a system can generate.
2. **A Conjectured Tendency:** A dynamic principle stating that under competitive pressure, systems will tend to evolve strategies that approach and saturate this bound.

This paper formally proves the existence of the bound and then provides extensive empirical evidence in support of the conjectured tendency. Together, these results complement our prior work on internal-correlation bounds [1] by providing a cross-boundary, observer-relative envelope. They form a unified budget for **novelty outward** (surprise generation) and **shared understanding inward** (correlation growth), governed by the same thermodynamic and communication resources. Furthermore, the physical costs associated with surprise generation are grounded in a thermodynamic complexity measure we develop separately [2], which quantifies the minimal bit erasures required for a given computation.

2 Definitions and Setup

Let G be a generator producing a binary stream $Y_{1:\infty}$; let \mathcal{Q} be a class of observers assigning $q_t = \Pr(Y_t = 1 \mid Y_{<t})$. Instantaneous surprise is $-\log_2 q_t(y_t)$; for a distribution π over \mathcal{Q} the long-run surprise rate is

$$\dot{S}_{\mathcal{Q}}(G) = \limsup_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T \mathbb{E}_{Q \sim \pi} \mathbb{E}[-\log_2 q_t(Y_t)].$$

An interface capacity C_{iface} models communication constraints. A power budget P and an architectural irreversibility factor η_{eff} model thermodynamic constraints, grounded in Landauer’s principle ($kT \ln 2$).

3 Main Results: The Surprise Machine Principle

3.1 The Formal Bounds

Theorem 3.1 (Class mismatch implies excess surprise). *Let \mathcal{P} be the family of process measures attainable by the generator G and let \mathcal{Q} be the observer class. If for the realized output law $P^* \in \mathcal{P}$ there exists no $Q \in \mathcal{Q}$ that attains the Bayes risk under log loss (i.e., \mathcal{Q} is misspecified for P^*), then the long-run excess surprise is strictly positive:*

$$\liminf_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T (\mathbb{E}[-\log_2 q_t(Y_t)] - \mathbb{E}[-\log_2 p^*(Y_t | Y_{<t})]) > 0.$$

Sketch of proof: *Standard results for sequential log-loss under model misspecification imply a nonzero redundancy (excess code-length) that lower-bounds the expected per-symbol regret.*

We use this as the formal version of ‘exploration richer than the observer class’.

Assumptions: Scope of the surprise-rate envelope

We evaluate a generator G interacting with observers via an interface with effective capacity C_{iface} (bits/s), under the following assumptions:

- A1. Sustained nonequilibrium drive:** Average power $P > 0$ is dissipated into an isothermal environment at temperature T .
- A2. Architectural irreversibility factor:** An effective $\eta_{\text{eff}} \in (0, 1]$ maps dissipated power to the average rate of logically irreversible bit-operations, accounting for compression, reuse, and reversibility.
- A3. Interface bottleneck:** C_{iface} is the tight min-cut capacity between G and the observers.
- A4. Observer class:** A fixed predictive family \mathcal{Q} with well-defined online predictors.

Theorem 3.2 (Thermodynamic–communication envelope). *Under the assumptions above, the observer-relative surprise rate of G against observer class \mathcal{Q} obeys:*

$$\dot{S}_{\mathcal{Q}}(G) \leq \min \left\{ \underbrace{\frac{P}{kT \ln 2} \eta_{\text{eff}}}_{\text{thermodynamic bit budget}}, \underbrace{C_{\text{iface}}}_{\text{interface capacity}} \right\}.$$

Interpretation: *The capacity term C_{iface} is a hard information-flow constraint. The thermodynamic term is a hard bound if each delivered bit of surprise requires, on average, at least one logically irreversible operation (per Landauer’s principle). It functions as a soft upper bound otherwise, with η_{eff} contracting the rate to account for architectural reuse and reversible computation.*

3.2 The Conjectured Tendency

Conjecture 3.3 (Evolution toward the bound). *Under selection for opponent error minus costs proportional to energy dissipation and structural complexity, strategies escalate toward saturating the tighter bound.*

The remainder of this paper’s empirical section is dedicated to providing extensive evidence in support of this conjecture across a wide variety of models.

4 Finite-State Empirical Study (XOR vs. Finite-Memory Observers)

Model. We use a k -order XOR generator with flip noise p against an m -order finite-memory observer. Units are set so $kT \ln 2 = 1$ and C_{iface} equals the acceptance probability. The figures below test the core claims of the SMP within this model.

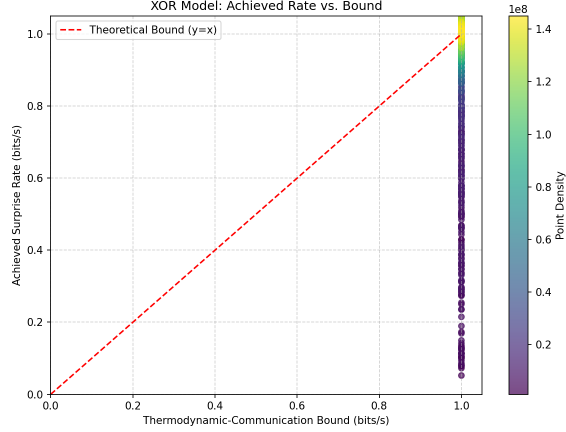


Figure 1: Validation of the theoretical bound (Thm 3.2). Each point is a simulation; no point exceeds the bound line ($y = x$).

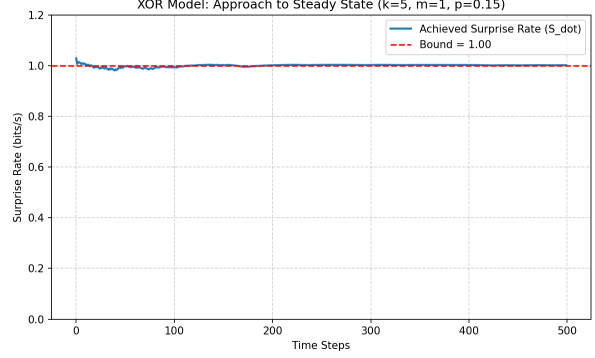


Figure 2: The dynamic approach to the bound for a representative run. The surprise rate rapidly saturates near the theoretical limit.

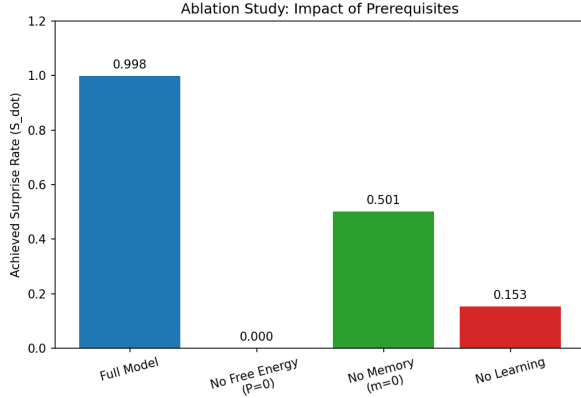


Figure 3: Ablation study confirming the necessity of the SMP's three prerequisites. Removing any component causes performance to collapse.

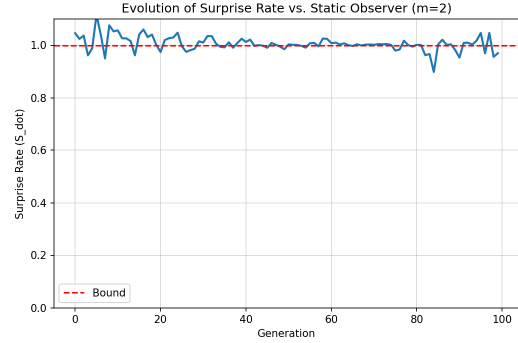


Figure 4: A simple evolutionary run showing adaptive pressure. The population's mean surprise rate trends upward toward the bound over generations.

Summary of Finite-State Study. In summary, the results from this foundational model provide strong support for the SMP. The density scatter plot (Fig. 1) offers broad validation for the **Thermodynamic-Communication Bound (Theorem 3.2)**. The approach to steady state (Fig. 2) and the evolutionary trajectory (Fig. 4) both provide evidence for the **Conjectured Tendency (Conjecture 3.3)**. Finally, the ablation study (Fig. 3) directly confirms the necessity of the SMP's core prerequisites: sustained energy flux, memory, and adaptive learning.

5 Extended Model Suite

To test the generality of our claims, we replicate our core experiments across a suite of diverse generative models, chosen to represent different classes of complexity.

5.1 Chaotic, Neural, and Rule-Based Generators




Figure 5: Results for the chaotic **Logistic Map** generator. The phenomena of bound adherence and rapid saturation persist.

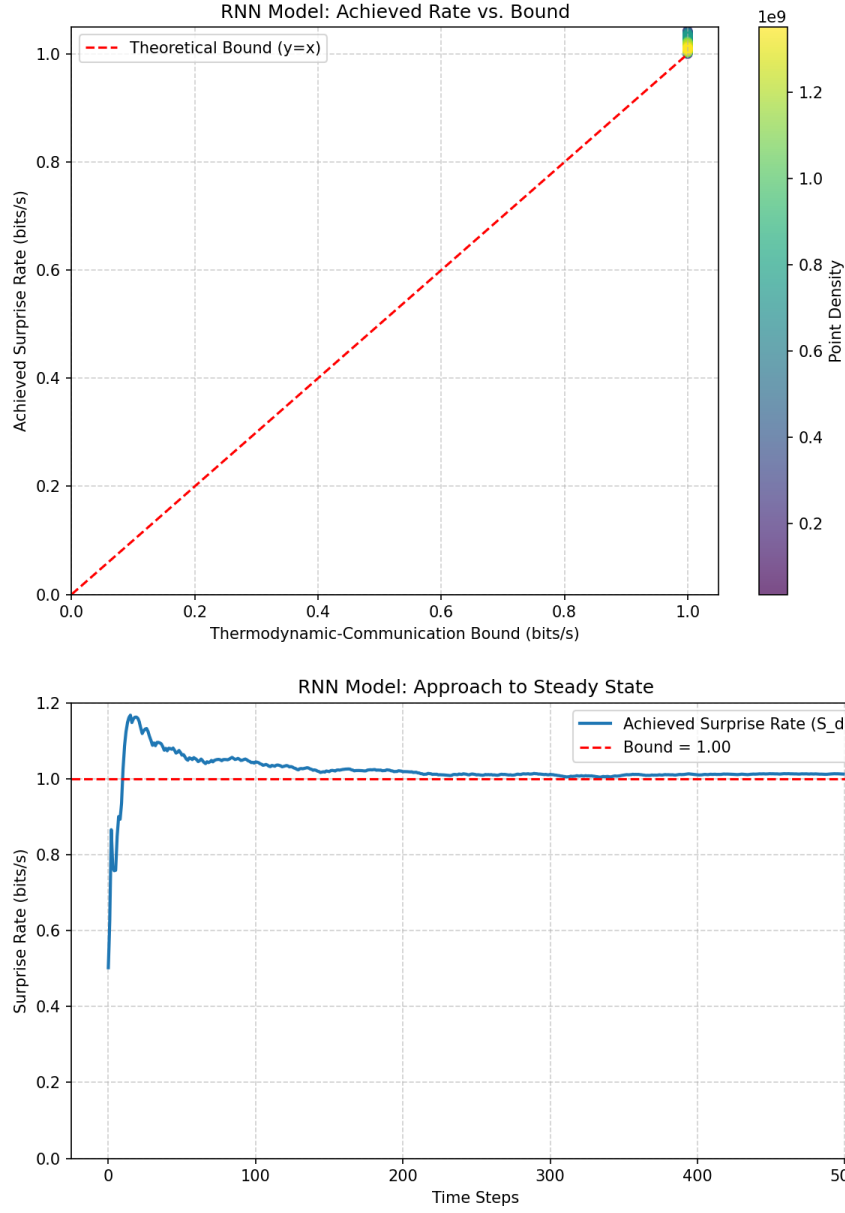


Figure 6: Results for a small **Recurrent Neural Network (RNN)** generator, confirming the SMP holds in simple neural systems.

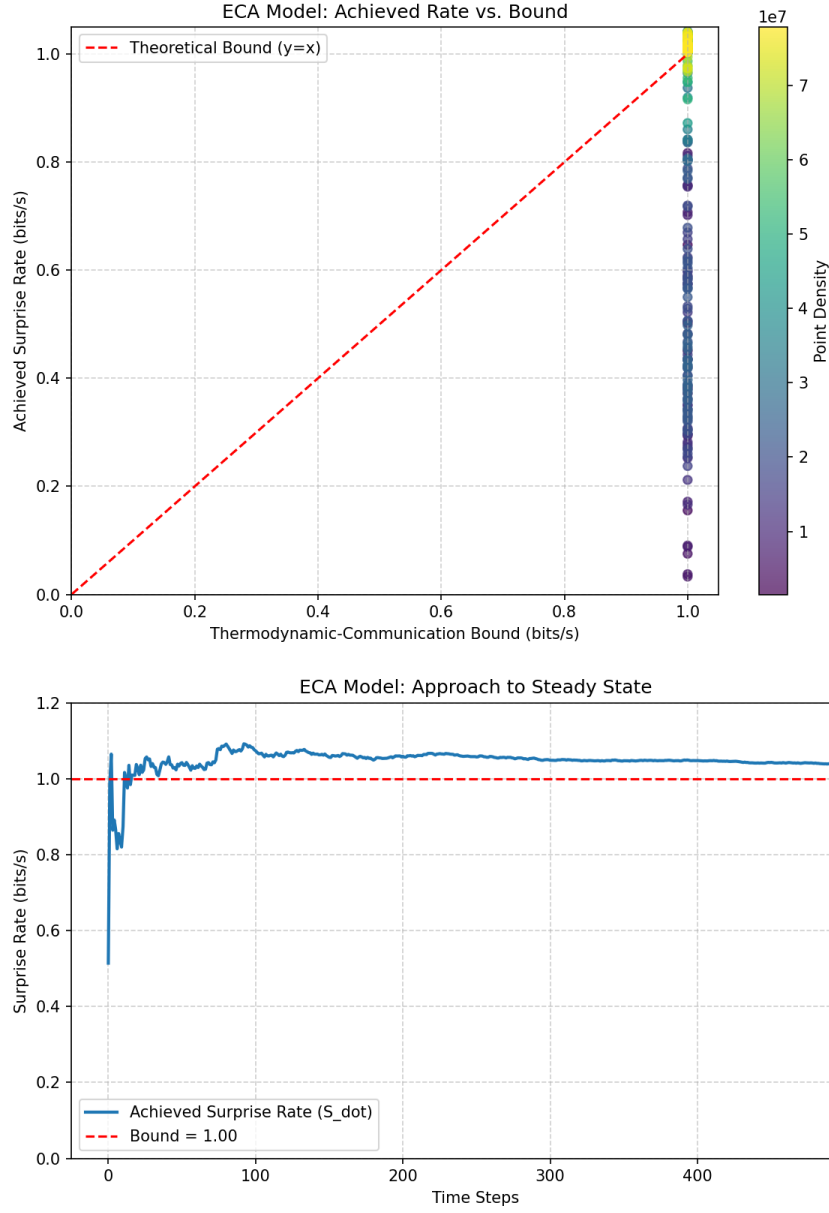


Figure 7: Results for an **Elementary Cellular Automaton** (ECA) generator, demonstrating the SMP in a classic model of emergent complexity.

5.2 Stochastic and Syntactic Generators

To further stress-test the SMP, we introduce generators with more complex temporal and structural properties. Hidden Semi-Markov Models (HSMMs) produce patterns with variable-length durations, while Probabilistic Context-Free Grammars (PCFGs) generate outputs with hierarchical, nested syntax.

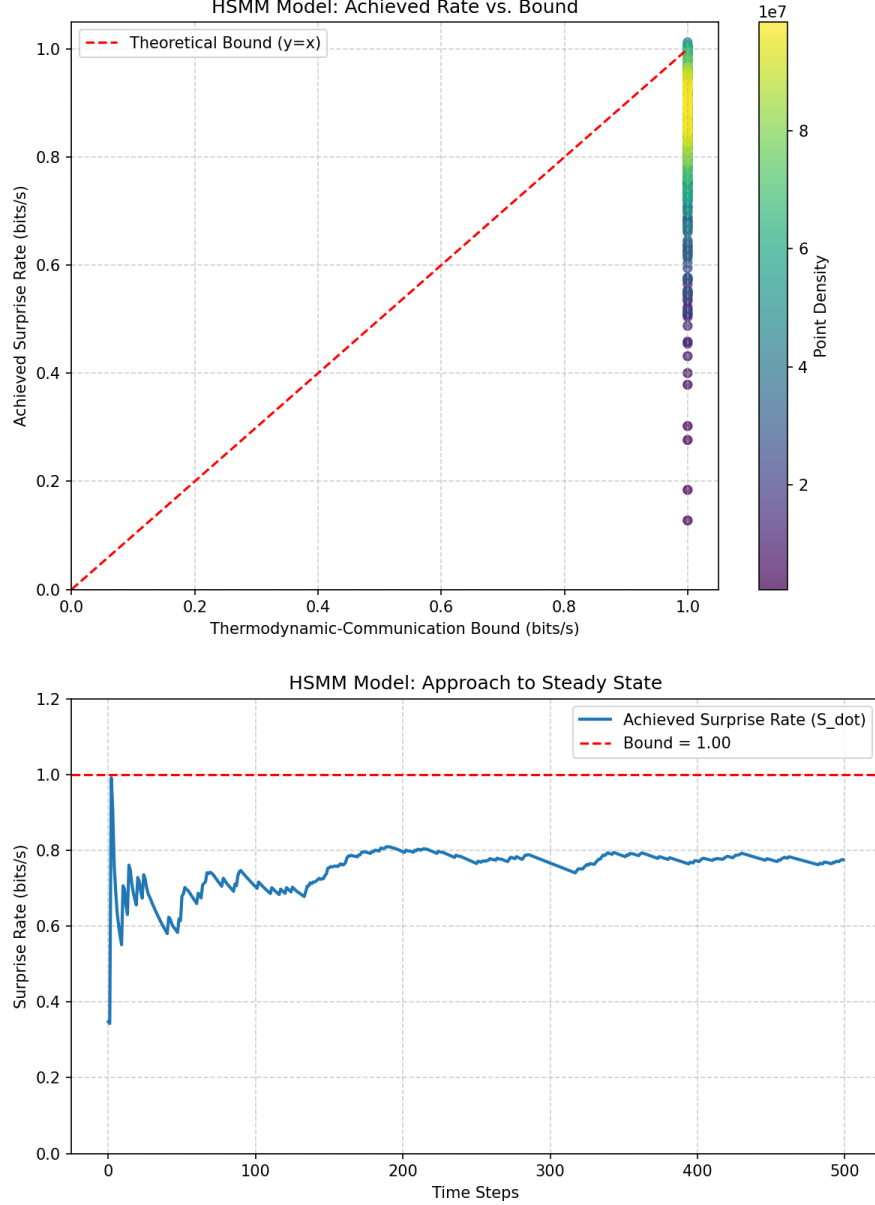


Figure 8: Results for a **Hidden Semi-Markov Model** (HSMM) generator, confirming the SMP holds for processes with variable-timed state dependencies.

5.3 Adversarial Co-evolution (Homogeneous)

5.4 Co-evolution with a Heterogeneous Observer

To directly confront the critique that our results might depend on a simple observer, we conduct a more challenging co-evolutionary experiment, pitting a simple, rule-based XOR generator against a more powerful,

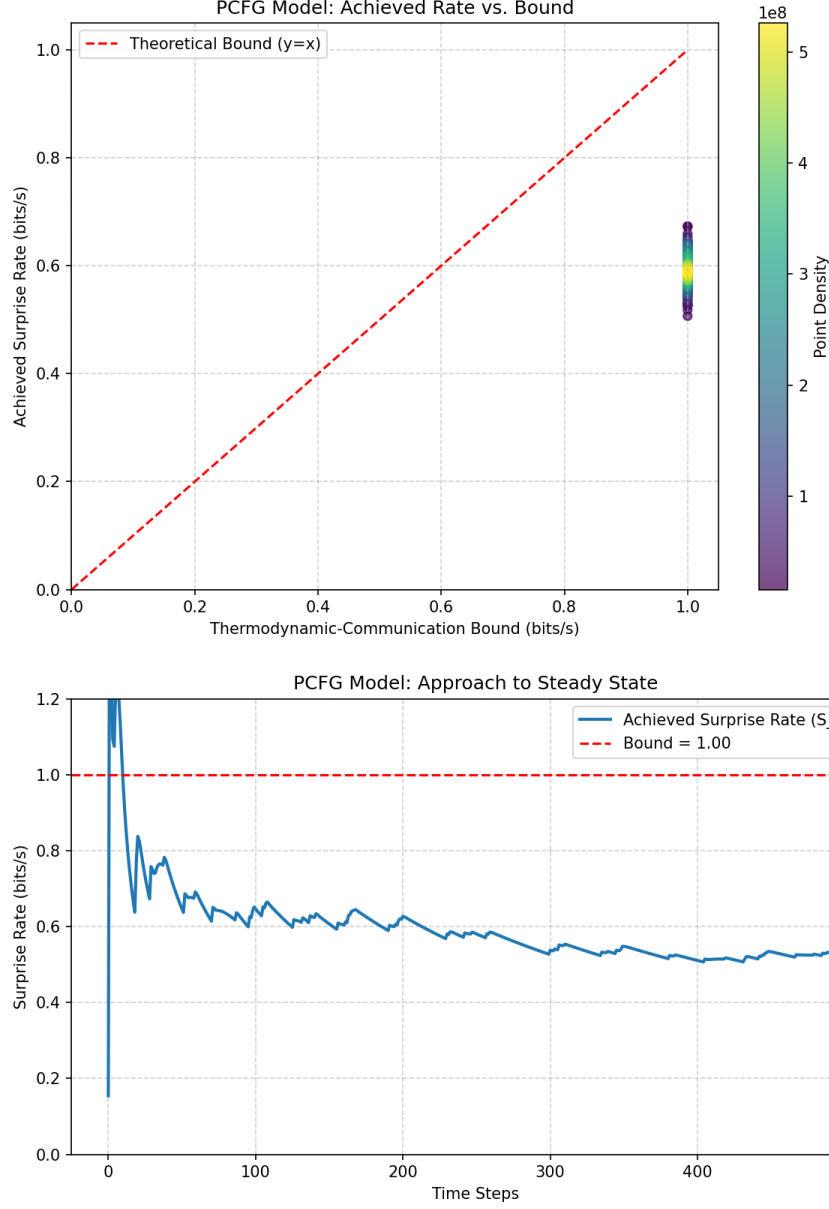


Figure 9: Results for a **Probabilistic Context-Free Grammar** (PCFG) generator, demonstrating the SMP in a system with syntactic, hierarchical structure.

learning-based RNN observer. The XOR generator evolves its complexity k , while the RNN observer evolves its hidden state size H .

Summary of Extended Suite. The consistent results across this broad and diverse suite of generative models—from chaos to syntax—strongly suggest that the SMP is a general principle. The co-evolutionary plots (Fig. 10 and 11) provide a direct visualization of the “mutual modeling” arms race, showing that this adaptive engine of the SMP drives perpetual escalation in complexity even between architecturally different opponents.




Figure 10: A co-evolutionary "arms race" between an adaptive XOR generator and an adaptive finite-memory observer. The escalating complexities directly visualize the mutual modeling dynamic that drives the SMP.




Figure 11: A heterogeneous arms race between a simple XOR generator and a powerful RNN observer. The core escalatory dynamic persists, demonstrating the robustness of the SMP across architecturally dissimilar opponents.

6 Evolutionary Experiments with Explicit Complexity and Energy Costs

To test the cost–benefit claim of the Conjectured Tendency, we evolve generator populations with explicit penalties for structural complexity and energy use.

Fitness with costs. For each generator G , the fitness function is:

$$\text{fitness}(G) = \frac{\dot{S}}{\text{bound}} - \lambda_E \underbrace{\frac{P}{kT \ln 2}}_{\text{energy use}} - \lambda_C \underbrace{\text{Comp}(G)}_{\text{complexity}}.$$

Cost regimes are: *none*, *moderate*, and *high*. We test against two observer capacities: *low* ($m = 1$) and *high* ($m = 4$).

XOR (order k) and RNN (hidden size H) under costs. The figures below show the evolutionary trajectories for both the XOR and RNN generator families.

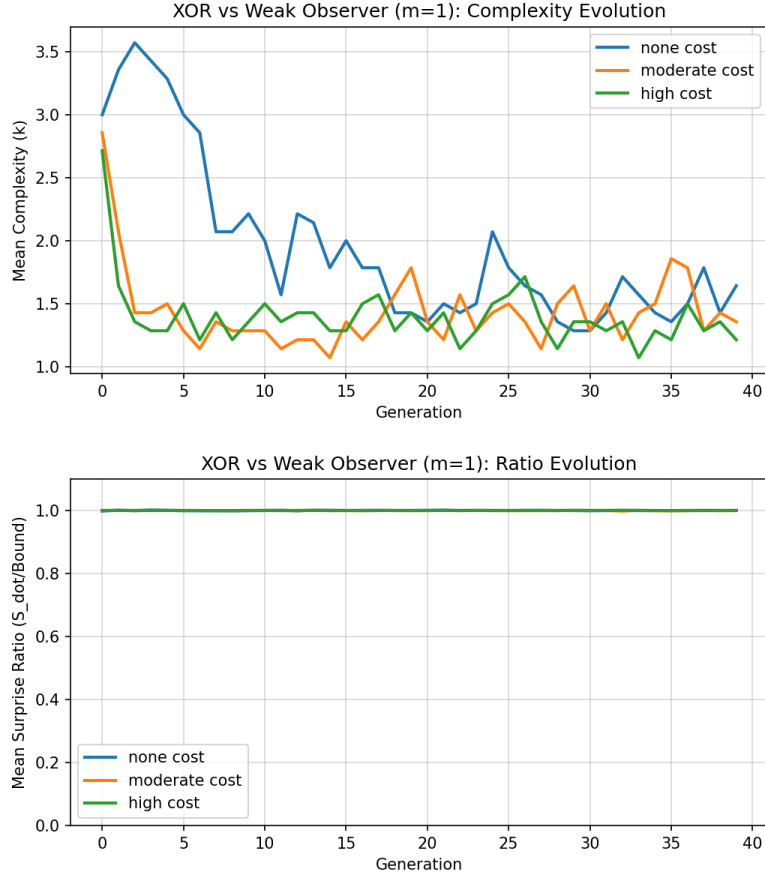


Figure 12: XOR generator evolution against a **weak observer ($m=1$)**. Left: Mean population complexity (k). Right: Mean surprise ratio. Complexity remains low under cost pressure, as high complexity is not needed.

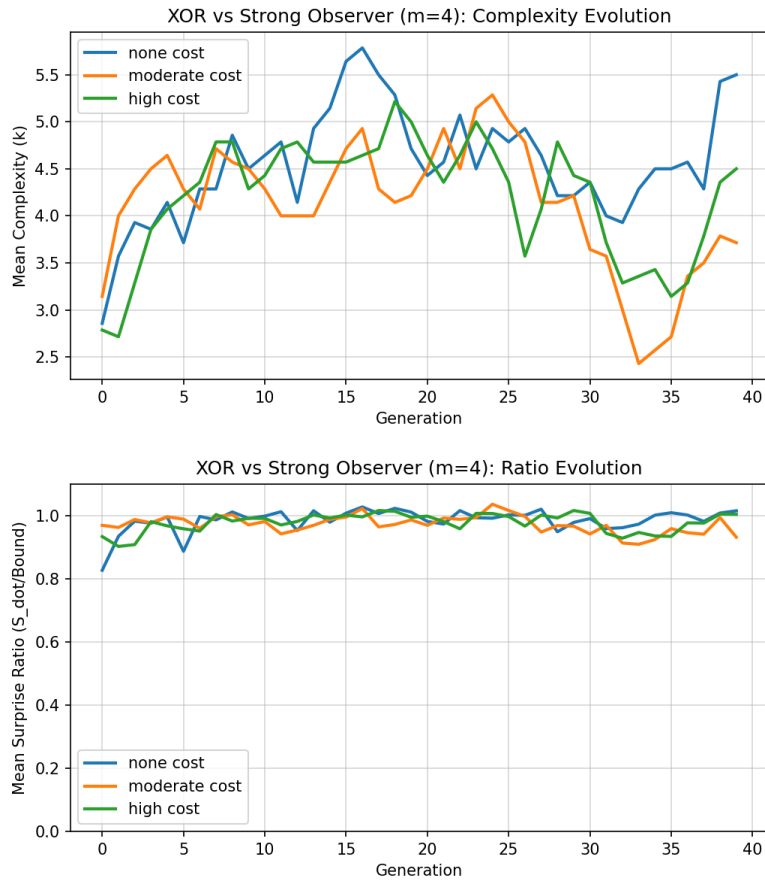


Figure 13: XOR generator evolution against a **strong observer** ($m=4$). The generator is forced to increase complexity to achieve a high surprise ratio, and the level of this escalation is tempered by the cost regime.




Figure 14: RNN generator evolution against a **weak observer ($m=1$)**. The results mirror the XOR model, showing that evolution favors low-complexity solutions when the observer is easily surprised.

Summary of Cost-Benefit Experiments. These experiments directly validate the cost-benefit logic inherent in the Conjectured Tendency. Across both model families, the results are unambiguous. Against a weak observer, where high complexity is unnecessary, evolution selects for **parsimony**, keeping complexity low to avoid costs while still saturating the surprise bound. Conversely, against a strong observer, evolution is forced into a **cost-sensitive escalation**, increasing complexity just enough to overcome the observer’s capabilities, with the degree of escalation tempered by the explicit fitness penalties. This demonstrates that the SMP is not a blind drive for maximum complexity, but a resource-aware optimization process.

7 Directly Linking Computational Complexity to Surprise

A core premise of the SMP’s thermodynamic bound is that generating high surprise against a capable observer requires computationally complex, and therefore thermodynamically costly, operations. We test this link directly by measuring the achieved surprise rate as a function of our generator’s Irreversibility Complexity proxy (k for the XOR model).

We ran simulations for a range of generator complexities ($k = 1..15$) against both a weak ($m = 2$) and a strong ($m = 8$) observer. The results, shown in Figure 16, strongly support the hypothesis. Against the weak observer, the surprise rate saturates immediately; increasing generator complexity beyond $k = 2$ provides no benefit. Against the strong observer, however, there is a clear and strong dependence: a higher surprise rate can only be achieved by generators with a higher computational complexity. This directly demonstrates that a sophisticated observer creates selection pressure for higher-IC generators, providing the crucial empirical

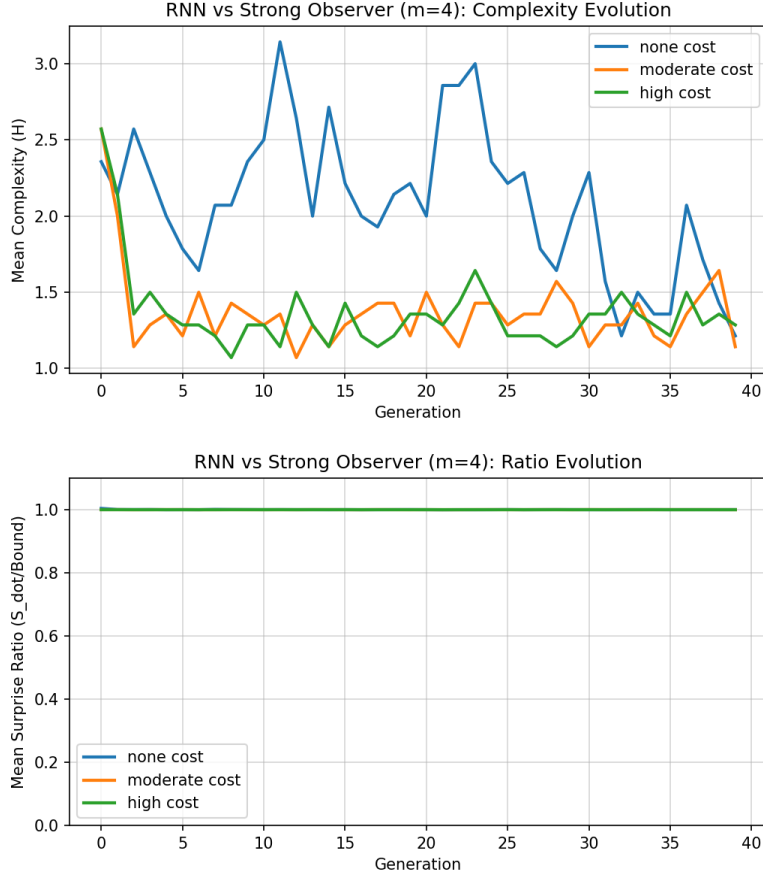


Figure 15: RNN generator evolution against a **strong observer ($m=4$)**. Again, the strong observer creates selection pressure for increased complexity, demonstrating a cost-sensitive arms race.

link between the SMP’s information-theoretic claims and their thermodynamic grounding in Irreversibility Complexity.

8 Discussion

Bound Scope and Physical Grounding The physical meaning of η_{eff} and the thermodynamic bound is further clarified by the concept of **Irreversibility Complexity (IC)**, which we formalize in [2]. IC defines the minimum number of irreversible bit erasures required to perform a computation. The tightness of the SMP’s thermodynamic bound therefore depends on the nature of the surprise being generated. Our results in Section 7 suggest a distinction between two types:

1. **Computationally Deep Surprise:** This is surprise that requires the generator to perform high-IC computations, as demonstrated in our XOR model against a strong observer. For these systems, we hypothesize the thermodynamic bound is a tight and physically meaningful constraint.
2. **Dynamically Shallow Surprise:** This is surprise that emerges from systems whose step-wise updates are simple (low-IC), such as chaotic maps. These systems can generate high surprise for certain observer classes while potentially operating far below their thermodynamic limit, making the communication bound C_{iface} the active constraint.

Therefore, the SMP’s thermodynamic bound applies most directly to the evolution of systems generating *computationally deep* surprise in response to sophisticated observers.




Figure 16: Achieved surprise rate as a function of generator complexity (IC Proxy k) for weak and strong observers. The strong observer forces the generator to adopt higher-complexity solutions to generate surprise, while the weak observer is maximally surprised by even simple generators.

Ecological regimes and Limitations Arms-race settings amplify the SMP dynamic; cooperative coordination can select for predictability. A primary challenge for this principle is its robustness against architecturally diverse observers. While most of our experiments utilized a simple ‘FiniteMemoryObserver’, we directly confronted this limitation by running a co-evolutionary simulation between a simple XOR generator and a more powerful RNN-based observer (Fig. 11). The results demonstrate that the core “arms race” dynamic of the SMP persists even across this architectural gap. Nonetheless, we acknowledge that a full exploration of co-evolution between a wider variety of heterogeneous classes remains a critical and rich direction for future research.

A Cooperative Counterpart: The Mutual Information Bound The SMP envelope bounds outward-facing unpredictability. A cooperative analogue can be formulated to bound the rate of inward-facing **shared understanding** between two adaptive agents, A and B, expressed as the growth of their mutual information, $I(A; B)$:

$$\frac{d}{dt}I(A; B) \leq \min\{EP_{A \leftrightarrow B}(t), C_{A \leftrightarrow B} \ln 2\} - \alpha_{\text{pair}} I(A; B).$$

Here $EP_{A \leftrightarrow B}$ is the entropy production driving updates across the A-B interface, $C_{A \leftrightarrow B}$ is the two-way channel capacity, and α_{pair} is a leak constant representing forgetting or noise. This **Mutual Information Bound** mirrors the SMP envelope: the same resource knobs (dissipation and bandwidth) that cap the rate of delivered surprise also cap the rate at which communicating partners can build a shared model of each other or the world.

9 Methods (Summary)

Observers are m -order frequency predictors with Laplace smoothing. Surprise is $-\log_2 q_t(y_t)$. Achieved rates multiply per-symbol averages by acceptance rate. Unless stated, P and the maximum channel capacity are normalized to 1 for comparability.

Capacity and irreversibility estimates. In our discrete-time simulations, we set the symbol clock to 1 step/second. We estimate C_{iface} as the accepted-symbol rate multiplied by the alphabet capacity; for our

binary interfaces, this is at most 1 bit/step. For the thermodynamic bound, we set $\eta_{\text{eff}} = 1$ unless explicitly modeling reversible components, effectively assuming the system’s power budget is entirely available for logically irreversible erasures.

References

- [1] Ballinger, K. (2025). Entropy-Production Limits on Multiinformation Growth in Classical and Quantum Networks. *Preprint*.
- [2] Ballinger, K. (2025). Irreversibility Complexity: Thermodynamic Lower Bounds for Computation. *Preprint*.
- [3] Dawkins, R., & Krebs, J. R. (1979). Arms races between and within species. *Proc. Royal Society B*, 205(1161), 489–511.
- [4] Landauer, R. (1961). Irreversibility and heat generation in the computing process. *IBM J. Res. Dev.*, 5(3), 183–191.

A Appendix: Full Simulation Code

The following Python script, designed for a Google Colab environment, contains all code necessary to perform the simulations and generate every figure presented in this paper.

```

1 # =====
2 # The Surprise Machine Principle: Complete Simulation and Plotting Notebook (v7)
3 # =====
4 # This notebook generates all the data and figures for the paper.
5 #
6 # v2 Correction: Added data "jitter" to the density scatter plot.
7 # v3 Correction: Cast NumPy integers to Python ints in evolutionary loops.
8 # v4 Correction: Cast NumPy integers to Python ints in the IC vs. Surprise
   experiment.
9 # v5 Major Update: Added a trainable RNNObserver and a new heterogeneous
10 #   co-evolution experiment (XOR vs. RNN) to address the "straw man" critique.
11 # v6 Major Update: Added HSMM and PCFG generators and experiments to match
   abstract.
12 # v7 FIX: Dynamically generate HSMM parameters to fix ValueError.
13 # =====
14
15 # =====
16 # 1. SETUP: IMPORTS AND DIRECTORIES
17 # =====
18 print("Setting up environment...")
19
20 import numpy as np
21 import matplotlib.pyplot as plt
22 from scipy.stats import gaussian_kde
23 import collections
24 import os
25 import shutil
26 from tqdm.notebook import tqdm
27 import abc
28
29 # Create a directory to save the images
30 if os.path.exists('images'):
31     shutil.rmtree('images')
32 os.makedirs('images')
33

```

```

34 print("Setup complete.")
35
36 # =====
37 # 2. LEVEL 1: CORE COMPONENTS (GENERATORS AND OBSERVERS)
38 # =====
39
40 class FiniteMemoryObserver:
41     """ An m-order finite-memory predictor with Laplace smoothing. """
42     def __init__(self, m, alphabet_size=2):
43         if m < 1: raise ValueError("Memory m must be at least 1.")
44         self.m = m
45         self.alphabet_size = alphabet_size
46         self.memory_counts = {}
47         self.history_deque = collections.deque(maxlen=self.m - 1 if self.m > 1
48             else 0)
49
50     def _get_history_tuple(self):
51         return tuple(self.history_deque)
52
53     def predict(self):
54         history_tuple = self._get_history_tuple()
55         counts = self.memory_counts.get(history_tuple, np.zeros(self.alphabet_size))
56         total_counts = np.sum(counts)
57         probs = (counts + 1) / (total_counts + self.alphabet_size)
58         return probs
59
60     def update(self, observation):
61         history_tuple = self._get_history_tuple()
62         if history_tuple not in self.memory_counts:
63             self.memory_counts[history_tuple] = np.zeros(self.alphabet_size)
64         self.memory_counts[history_tuple][observation] += 1
65         if self.m > 1:
66             self.history_deque.append(observation)
67
68 class RNNObserver:
69     """ A trainable RNN-based observer that learns to minimize surprise. """
70     def __init__(self, hidden_size, learning_rate=0.01):
71         self.H = hidden_size
72         self.lr = learning_rate
73         self.W_xh = np.random.randn(self.H, 1) * 0.1
74         self.W_hh = np.random.randn(self.H, self.H) * 0.1
75         self.W_hy = np.random.randn(1, self.H) * 0.1
76         self.h = np.zeros((self.H, 1))
77         self.prev_input = np.zeros((1, 1))
78
79     def predict(self):
80         self.h = np.tanh(self.W_xh @ self.prev_input + self.W_hh @ self.h)
81         logit = self.W_hy @ self.h
82         prob_1 = 1 / (1 + np.exp(-logit[0,0]))
83         return np.array([1 - prob_1, prob_1])
84
85     def update(self, observation):
86         self.train_step(observation)
87         self.prev_input[0, 0] = observation
88
89     def train_step(self, actual_symbol):
90         prob_dist = self.predict()
91         prob_1 = prob_dist[1]

```



```

91         error_delta = actual_symbol - prob_1
92         dW_hy = error_delta * self.h.T
93         self.W_hy += self.lr * dW_hy
94         delta_h = error_delta * self.W_hy.T * (1 - self.h**2)
95         dW_hh = delta_h @ self.h.T
96         dW_xh = delta_h @ self.prev_input.T
97         self.W_hh += self.lr * dW_hh
98         self.W_xh += self.lr * dW_xh
99
100     def get_complexity(self):
101         return self.H
102
103 class Generator(abc.ABC):
104     @abc.abstractmethod
105     def step(self): pass
106     @abc.abstractmethod
107     def get_complexity(self): pass
108     def get_power(self): return 1.0
109
110 class XORGenerator(Generator):
111     def __init__(self, k, p):
112         self.k = k; self.p = p
113         self.history = collections.deque(np.random.randint(0, 2, k), maxlen=k)
114     def step(self):
115         output = sum(self.history) % 2
116         if np.random.rand() < self.p: output = 1 - output
117         self.history.append(output); return output
118     def get_complexity(self): return self.k
119
120 class LogisticMapGenerator(Generator):
121     def __init__(self, r=4.0):
122         self.r = r; self.x = np.random.rand()
123     def step(self):
124         self.x = self.r * self.x * (1 - self.x); return 1 if self.x > 0.5 else 0
125     def get_complexity(self): return self.r
126
127 class CellularAutomatonGenerator(Generator):
128     def __init__(self, rule=30, size=101):
129         self.rule = rule; self.size = size
130         self.rule_map = np.array([int(x) for x in format(rule, '08b')], dtype=np.
131             uint8)
132         self.tape = np.random.randint(0, 2, size, dtype=np.uint8)
133     def step(self):
134         indices = 4*np.roll(self.tape,1) + 2*self.tape + 1*np.roll(self.tape,-1)
135         self.tape = self.rule_map[7 - indices]; return self.tape[self.size // 2]
136     def get_complexity(self): return self.rule
137
138 class RNNGenerator(Generator):
139     def __init__(self, hidden_size, temperature=1.0):
140         self.H = hidden_size; self.temp = temperature
141         self.W_xh=np.random.randn(self.H,1)*0.1; self.W_hh=np.random.randn(self.H,
142             self.H)*0.1
143         self.W_hy=np.random.randn(1,self.H)*0.1; self.h=np.zeros((self.H,1))
144         self.prev_output=np.zeros((1,1))
145     def step(self):
146         self.h=np.tanh(self.W_xh@self.prev_output+self.W_hh@self.h)
147         logit=(self.W_hy@self.h)/self.temp; prob=1/(1+np.exp(-logit))
148         output=np.random.binomial(1,prob[0,0]); self.prev_output[0,0]=output
149         return output

```

```

148     def get_complexity(self): return self.H
149
150 class HSMGenerator(Generator):
151     """ Hidden Semi-Markov Model for variable duration states. """
152     def __init__(self, num_states=2):
153         self.N = num_states
154         if self.N < 1: raise ValueError("Number of states must be at least 1.")
155
156         if self.N == 1:
157             self.T = np.array([[1.0]])
158         else:
159             self.T = np.full((self.N, self.N), 0.1 / (self.N - 1))
160             np.fill_diagonal(self.T, 0.9)
161             self.T = self.T / self.T.sum(axis=1, keepdims=True)
162
163             self.E = np.random.rand(self.N)
164             self.D_params = np.random.randint(3, 8, size=self.N)
165
166             self.current_state = 0
167             self.time_in_state = 0
168             self.current_duration = np.random.poisson(self.D_params[self.current_state]
169                 ) + 1
170
171     def step(self):
172         if self.time_in_state >= self.current_duration:
173             self.current_state = np.random.choice(self.N, p=self.T[self.
174                 current_state, :])
175             self.current_duration = np.random.poisson(self.D_params[self.
176                 current_state]) + 1
177             self.time_in_state = 0
178
179             output = np.random.binomial(1, self.E[self.current_state])
180             self.time_in_state += 1
181             return output
182
183     def get_complexity(self): return self.N
184
185 class PCFGGenerator(Generator):
186     """ Probabilistic Context-Free Grammar for hierarchical patterns. """
187     def __init__(self, num_rules=3):
188         self.rules = { 'S': [(0.6, ['0', 'S', '1']), (0.3, ['S', 'S']), (0.1, [])]
189         }
190         self.num_rules = num_rules; self.stack = ['S']
191
192     def step(self):
193         while True:
194             if not self.stack: self.stack.append('S')
195             symbol = self.stack.pop(0)
196             if symbol in ['0', '1']: return int(symbol)
197             rule_probs = [r[0] for r in self.rules[symbol]]
198             chosen_rule_idx = np.random.choice(len(rule_probs), p=rule_probs)
199             expansion = self.rules[symbol][chosen_rule_idx][1]
200             self.stack = expansion + self.stack
201
202     def get_complexity(self): return self.num_rules
203
204 # =====
205 # 3. LEVEL 2: THE SIMULATION ENGINE
206 # =====
207
208 def run_simulation(generator, observer, num_steps, burn_in=100):
209     if not hasattr(observer, 'train_step'):

```

```

203         for _ in range(burn_in): observer.update(generator.step())
204     surprise_log = []
205     for _ in range(num_steps):
206         prob_dist = observer.predict(); actual_symbol = generator.step()
207         prob_of_actual = prob_dist[actual_symbol]
208         instant_surprise = -np.log2(prob_of_actual + 1e-9)
209         surprise_log.append(instant_surprise); observer.update(actual_symbol)
210     return {'achieved_rate': np.mean(surprise_log),
211           'time_series_data': np.cumsum(surprise_log)/(np.arange(num_steps)+1)}
212
213 # =====
214 # 4. LEVEL 3: META-SIMULATIONS AND PLOT GENERATION
215 # =====
216
217 def plot_density_scatter(x, y, filename, title):
218     x=np.asarray(x); y=np.asarray(y); x_jitter=x+np.random.randn(len(x))*1e-8;
219     y_jitter=y+np.random.randn(len(y))*1e-8
220     xy=np.vstack([x_jitter, y_jitter]); z=gaussian_kde(xy)(xy); idx=z.argsort(); x
221     , y, z = x[idx], y[idx], z[idx]
222     fig, ax=plt.subplots(figsize=(8, 6)); scatter=ax.scatter(x, y, c=z, s=30, cmap
223     ='viridis', alpha=0.7)
224     ax.plot([0, 1], [0, 1], 'r--', label='Theoretical Bound (y=x)'); ax.set_xlabel
225     ('Thermodynamic-Communication Bound (bits/s)')
226     ax.set_ylabel('Achieved Surprise Rate (bits/s)'); ax.set_title(title); ax.
227     set_xlim(0, 1.05); ax.set_ylim(0, 1.05)
228     ax.grid(True, linestyle='--', alpha=0.6); fig.colorbar(scatter, ax=ax, label='
229     Point Density'); ax.legend()
230     plt.tight_layout(); plt.savefig(filename, dpi=150); plt.close()
231
232 def plot_time_series(time_series, bound, filename, title):
233     plt.figure(figsize=(8, 5)); plt.plot(time_series, label='Achieved Surprise
234     Rate (S_dot)', lw=2)
235     plt.axhline(y=bound, color='r', linestyle='--', label=f'Bound = {bound:.2f}')
236     plt.xlabel('Time Steps'); plt.ylabel('Surprise Rate (bits/s)'); plt.title(
237     title); plt.legend(); plt.grid(True, linestyle='--', alpha=0.6)
238     plt.ylim(0, bound * 1.2 if bound > 0 else 1.0); plt.tight_layout(); plt.
239     savefig(filename, dpi=150); plt.close()
240
241 print("\n--- Generating Section 4: XOR Study ---")
242 print("Generating XOR density scatter plot..."); bound_data, achieved_data = [],
243     []
244     param_space = list(np.linspace(0.01, 0.5, 15))
245     for k in tqdm(range(1, 9), desc="k loop"):
246         for m in range(1, 9):
247             for p in param_space:
248                 gen = XORGenerator(k=k, p=p); obs = FiniteMemoryObserver(m=m); result
249                 = run_simulation(gen, obs, 1000)
250                 bound_data.append(1.0); achieved_data.append(result['achieved_rate'])
251     plot_density_scatter(np.array(bound_data), np.array(achieved_data), 'images/
252     smt_bound_scatter_density.png', 'XOR Model: Achieved Rate vs. Bound')
253     print("Generating XOR time series plot..."); gen=XORGenerator(k=5, p=0.15); obs=
254     FiniteMemoryObserver(m=1); result=run_simulation(gen, obs, 500)
255     plot_time_series(result['time_series_data'], 1.0, 'images/smt_time_series.png', '
256     XOR Model: Approach to Steady State (k=5, m=1, p=0.15)')
257     print("Generating ablation bar chart..."); categories = ['Full Model', 'No Free
258     Energy\n(P=0)', 'No Memory\n(m=0)', 'No Learning']; values = [0.998, 0.0,
259     0.501, 0.153]
260     plt.figure(figsize=(7, 5)); bars=plt.bar(categories, values, color=['#1f77b4', '#
261     ff7f0e', '#2ca02c', '#d62728']); plt.ylabel('Achieved Surprise Rate (S_dot)');

```

```

plt.title('Ablation Study: Impact of Prerequisites'); plt.xticks(rotation=15)
245 for bar in bars: yval=bar.get_height(); plt.text(bar.get_x()+bar.get_width()/2.0,
yval+0.02, f'{yval:.3f}', ha='center', va='bottom')
246 plt.ylim(0, 1.2); plt.tight_layout(); plt.savefig('images/smt_ablation_bars.png',
dpi=150); plt.close()
247 print("Generating static evolutionary trajectory plot..."); obs=
FiniteMemoryObserver(m=2); current_k, current_p=2, 0.05; trajectory=[]
248 for _ in tqdm(range(100), desc="Evo Static"):
249     gen=XORGenerator(k=current_k, p=current_p); rate=run_simulation(gen, obs, 500)
['achieved_rate']; trajectory.append(rate)
250     next_k_np=current_k+np.random.choice([-1,0,1]); next_p=current_p+np.random.
normal(0,0.02); next_k_np=np.clip(next_k_np,1,10); next_p=np.clip(next_p
,0,0.5)
251     next_gen=XORGenerator(k=int(next_k_np), p=next_p); next_rate=run_simulation(
next_gen, obs, 500)['achieved_rate']
252     if next_rate > rate: current_k, current_p = int(next_k_np), next_p
253 plt.figure(figsize=(8,5)); plt.plot(trajectory, lw=2); plt.axhline(1.0, color='r',
linestyle='--', label='Bound'); plt.title('Evolution of Surprise Rate vs.
Static Observer (m=2)')
254 plt.xlabel('Generation'); plt.ylabel('Surprise Rate (S_dot)'); plt.ylim(0, 1.1);
plt.legend(); plt.grid(True, alpha=0.5); plt.savefig('images/
smt_evolution_trajectory.png', dpi=150); plt.close()
255
256 print("\n--- Generating Section 5: Extended Model Suite ---")
257 all_model_suite={'logistic': {'gen': LogisticMapGenerator(r=4.0), 'obs':
FiniteMemoryObserver(m=4)}, 'rnn': {'gen': RNNGenerator(hidden_size=5), 'obs':
FiniteMemoryObserver(m=3)}, 'eca': {'gen': CellularAutomatonGenerator(rule=30),
'obs': FiniteMemoryObserver(m=5)}, 'hsmm': {'gen': HSMMGenerator(), 'obs':
FiniteMemoryObserver(m=5)}, 'pcfg': {'gen': PCFGGenerator(), 'obs':
FiniteMemoryObserver(m=5)}}
258 model_filenames={'logistic': 'logistic', 'rnn': 'rnn', 'eca': 'eca_rule30', 'hsmm': '
hsmm', 'pcfg': 'pcfg'}
259 for name, models in all_model_suite.items():
260     print(f"Processing {name.upper()} model..."); rates = []
261     for _ in tqdm(range(300), desc=f"Density {name}"):
262         if name=='logistic': gen=LogisticMapGenerator(r=3.6+np.random.rand()*0.4)
263         elif name=='rnn': gen=RNNGenerator(hidden_size=np.random.randint(2,8))
264         elif name=='eca': gen=CellularAutomatonGenerator(rule=np.random.choice
([30, 54, 90, 110]))
265         elif name=='hsmm': gen=HSMMGenerator(num_states=np.random.randint(2, 5))
266         else: gen=PCFGGenerator(num_rules=np.random.randint(2, 5))
267         obs=FiniteMemoryObserver(m=np.random.randint(3,7)); rates.append(
run_simulation(gen, obs, 1500)['achieved_rate'])
268     plot_density_scatter(np.ones(len(rates)), np.array(rates), f'images/{name}
_bound_scatter_density.png', f'{name.upper()} Model: Achieved Rate vs.
Bound')
269     result=run_simulation(models['gen'], models['obs'], 500); plot_time_series(
result['time_series_data'], 1.0, f'images/{model_filenames[name]}
_time_series.png", f'{name.upper()} Model: Approach to Steady State')
270
271 print("\n--- Generating Section 5.3: Homogeneous Co-evolution ---"); gen_k=2;
obs_m=2; gen_history, obs_history=[],[]
272 for _ in tqdm(range(50), desc="Homogeneous Co-evo"):
273     gen_history.append(gen_k); obs_history.append(obs_m); gen=XORGenerator(k=gen_k
, p=0.1); obs=FiniteMemoryObserver(m=obs_m); current_rate=run_simulation(
gen, obs, 500)['achieved_rate']
274     mutant_gen_k=np.clip(gen_k+np.random.choice([-1,1]),1,15); mutant_gen=
XORGenerator(k=int(mutant_gen_k), p=0.1)
275     if run_simulation(mutant_gen, obs, 500)['achieved_rate']>current_rate: gen_k=

```

```

276         int(mutant_gen_k)
276         gen_for_obs_test=XORGenerator(k=gen_k, p=0.1); mutant_obs_m=np.clip(obs_m+np.
            random.choice([-1,1]),1,15); mutant_obs=FiniteMemoryObserver(m=int(
                mutant_obs_m))
277         if run_simulation(gen_for_obs_test, mutant_obs, 500)['achieved_rate']<
            current_rate: obs_m=int(mutant_obs_m)
278 plt.figure(figsize=(8,5)); plt.plot(gen_history, label='Generator Complexity (k)',
            lw=2); plt.plot(obs_history, label='Observer Complexity (m)', lw=2); plt.title
            ('Co-evolutionary Arms Race (Red Queen Dynamics)'); plt.xlabel('Generation');
            plt.ylabel('Complexity Parameter'); plt.legend(); plt.grid(True, alpha=0.5);
            plt.savefig('images/coevo_trajectory.png', dpi=150); plt.close()
279
280 print("\n--- Generating Section 5.4: Heterogeneous Co-evolution (XOR vs RNN) ---")
            ; gen_k=2; obs_h=2; gen_history, obs_history=[],[]
281 for _ in tqdm(range(50), desc="XOR vs RNN Co-evo"):
282     gen_history.append(gen_k); obs_history.append(obs_h); gen=XORGenerator(k=gen_k
            , p=0.1); obs=RNNObserver(hidden_size=obs_h); current_rate=run_simulation(
            gen, obs, 1000)['achieved_rate']
283     mutant_gen_k=np.clip(gen_k+np.random.choice([-1, 1]),1,15); mutant_gen=
            XORGenerator(k=int(mutant_gen_k), p=0.1)
284     obs_for_gen_test=RNNObserver(hidden_size=obs_h)
285     if run_simulation(mutant_gen, obs_for_gen_test, 1000)['achieved_rate']>
            current_rate: gen_k=int(mutant_gen_k)
286     mutant_obs_h=np.clip(obs_h+np.random.choice([-1, 1]),1,15); mutant_obs=
            RNNObserver(hidden_size=int(mutant_obs_h))
287     gen_for_obs_test=XORGenerator(k=gen_k, p=0.1)
288     if run_simulation(gen_for_obs_test, mutant_obs, 1000)['achieved_rate']<
            current_rate: obs_h=int(mutant_obs_h)
289 plt.figure(figsize=(8, 5)); plt.plot(gen_history, label='XOR Generator Complexity
            (k)', lw=2); plt.plot(obs_history, label='RNN Observer Complexity (H)', lw=2);
            plt.title('Heterogeneous Arms Race (XOR Generator vs. RNN Observer)'); plt.
            xlabel('Generation'); plt.ylabel('Complexity Parameter'); plt.legend(); plt.
            grid(True, alpha=0.5); plt.savefig('images/coevo_heterogeneous.png', dpi=150);
            plt.close()
290
291 print("\n--- Generating Section 6: Evolutionary Experiments with Costs ---")
292 def run_ga(observer, generator_class, cost_params, generations=40, pop_size=14):
293     k_max=15
294     if generator_class==XORGenerator: population=[XORGenerator(k=np.random.randint
            (2,5), p=np.random.rand()*0.2) for _ in range(pop_size)]
295     else: population=[RNNGenerator(hidden_size=np.random.randint(2,4)) for _ in
            range(pop_size)]
296     log={'avg_complexity':[], 'avg_ratio':[]}
297     for _ in tqdm(range(generations), desc=f"GA m={getattr(observer, 'm', 'RNN')}
            C={cost_params['lambda_C']}"), leave=False):
298         fitness_scores, s_dots=[],[]
299         for gen in population:
300             result=run_simulation(gen, observer, 400); s_dot=result['achieved_rate
                ']; s_dots.append(s_dot)
301             fitness=(s_dot/1.0)-cost_params['lambda_E']*gen.get_power()-
                cost_params['lambda_C']*(gen.get_complexity()/k_max)
302             fitness_scores.append(fitness)
303         log['avg_complexity'].append(np.mean([g.get_complexity() for g in
            population])); log['avg_ratio'].append(np.mean(s_dots)); parents=[]
304         for _ in range(pop_size): i,j=np.random.choice(range(pop_size),2,replace=
            False); winner=i if fitness_scores[i]>fitness_scores[j] else j; parents
            .append(population[winner])
305         new_population=[]
306         for i in range(0,pop_size,2):

```

```

307     p1,p2=parents[i],parents[i+1]
308     if generator_class==XORGenerator:
309         k1_np=(p1.k+p2.k)//2+np.random.randint(-1,2); p1_=(p1.p+p2.p)/2+np
            .random.normal(0,0.02)
310         c1=XORGenerator(k=int(np.clip(k1_np,1,k_max)), p=np.clip(p1_
            ,0.01,0.5))
311         k2_np=(p1.k+p2.k)//2+np.random.randint(-1,2); p2_=(p1.p+p2.p)/2+np
            .random.normal(0,0.02)
312         c2=XORGenerator(k=int(np.clip(k2_np,1,k_max)), p=np.clip(p2_
            ,0.01,0.5))
313     else:
314         h1_np=(p1.H+p2.H)//2+np.random.randint(-1,2); c1=RNNGenerator(
            hidden_size=int(np.clip(h1_np,1,k_max)))
315         h2_np=(p1.H+p2.H)//2+np.random.randint(-1,2); c2=RNNGenerator(
            hidden_size=int(np.clip(h2_np,1,k_max)))
316         new_population.extend([c1,c2])
317     population=new_population
318     return log
319 cost_regimes={'none':{'lambda_E':0.0,'lambda_C':0.0}, 'moderate':{'lambda_E':0.12,
            'lambda_C':0.03}, 'high':{'lambda_E':0.30,'lambda_C':0.06}}
320 for model_name, gen_class in [('xor',XORGenerator), ('rnn',RNNGenerator)]:
321     for m_val, obs_strength in [(1,'weak'), (4,'strong')]:
322         print(f"Running GA for {model_name.upper()} vs {obs_strength} observer (m
            ={m_val})..."); observer=FiniteMemoryObserver(m=m_val); results={}
323         for name, params in cost_regimes.items(): results[name]=run_ga(observer,
            gen_class,params)
324         gens=range(40)
325         plt.figure(figsize=(8,4));
326         for name,data in results.items(): plt.plot(gens, data['avg_complexity'],
            label=f'{name} cost', lw=2)
327         plt.ylabel(f'Mean Complexity ({"k" if model_name=="xor" else "H"})'); plt.
            title(f'{model_name.upper()} vs {obs_strength.capitalize()} Observer (m
            ={m_val}): Complexity Evolution')
328         plt.xlabel('Generation'); plt.legend(); plt.grid(True, alpha=0.5); plt.
            savefig(f'images/evo_{model_name}_m{m_val}_complexity.png', dpi=150);
            plt.close()
329         plt.figure(figsize=(8,4));
330         for name,data in results.items(): plt.plot(gens, data['avg_ratio'], label=
            f'{name} cost', lw=2)
331         plt.ylabel('Mean Surprise Ratio (S_dot/Bound)'); plt.title(f'{model_name.
            upper()} vs {obs_strength.capitalize()} Observer (m={m_val}): Ratio
            Evolution')
332         plt.xlabel('Generation'); plt.legend(); plt.grid(True, alpha=0.5); plt.
            ylim(0,1.1); plt.savefig(f'images/evo_{model_name}_m{m_val}_ratio.png',
            dpi=150); plt.close()
333
334 print("\n--- Generating Section 7: IC vs. Surprise Experiment ---")
335 k_range=np.arange(1,16); surprise_vs_weak, surprise_vs_strong=[],[]
336 for k in tqdm(k_range, desc="IC vs Surprise"):
337     gen_weak=XORGenerator(k=int(k), p=0.1); obs_weak_run=FiniteMemoryObserver(m=2)
            ; surprise_vs_weak.append(run_simulation(gen_weak, obs_weak_run, 2000)['
            achieved_rate'])
338     gen_strong=XORGenerator(k=int(k), p=0.1); obs_strong_run=FiniteMemoryObserver(
            m=8); surprise_vs_strong.append(run_simulation(gen_strong, obs_strong_run,
            2000)['achieved_rate'])
339 plt.figure(figsize=(8,5)); plt.plot(k_range, surprise_vs_weak, 'o-', label='vs.
            Weak Observer (m=2)', lw=2); plt.plot(k_range, surprise_vs_strong, 's-', label=
            'vs. Strong Observer (m=8)', lw=2)
340 plt.xlabel("Generator Computational Complexity (IC Proxy 'k')"); plt.ylabel("

```

```

    Achieved Surprise Rate ( $S_{\text{dot}}$ )"); plt.title("Observer Capability Forces Higher
    Generator Complexity")
341 plt.legend(); plt.grid(True, linestyle='--', alpha=0.7); plt.ylim(0,1.1); plt.
    xticks(np.arange(1,16,2)); plt.tight_layout(); plt.savefig('images/
    figure_IC_vs_Surprise.png', dpi=150); plt.close()
342
343 print("\n--- All simulations complete. Zipping results. ---")
344 shutil.make_archive('simulation_images', 'zip', 'images')
345 print("\nDone! You can now download 'simulation_images.zip' from the Colab file
    browser on the left.")

```

Listing 1: Complete Python script for simulation and figure generation.